

monte-carlo

November 3, 2017

```
In [1]: import numpy as np
import sklearn as skl
from matplotlib import pyplot as plt
%pylab inline
```

Populating the interactive namespace from numpy and matplotlib

1 1.1 Inverse Transform sampling

For $Y = 0$:

$$F_{Y=0} = \int_{-\inf}^x p_{Y=0}(x') dx' = \int_0^x 2 - 2x' dx' = 2x' - x'^2 \Big|_0^x \quad (1)$$

$$= 2x - x^2 \stackrel{!}{=} u, \quad u \text{ sampled from Unif}(0, 1) \quad (2)$$

$$\Rightarrow x^2 - 2x + u = 0 \quad (3)$$

$$\Rightarrow x = 1 \pm \sqrt{1 - u} \quad (4)$$

$$(5)$$

$$X \in [0, 1] \Rightarrow x_{Y=0} = 1 - \sqrt{1 - u} \quad (6)$$

For $Y = 1$:

$$F_{Y=1} = \int_{-\inf}^x p_{Y=1}(x') dx' = \int_0^x 2x' dx' = x'^2 \Big|_0^x \quad (7)$$

$$= x^2 \stackrel{!}{=} u, \quad u \text{ sampled from Unif}(0, 1) \quad (8)$$

$$\Rightarrow x = \pm \sqrt{u} \quad (9)$$

$$(10)$$

$$X \in [0, 1] \Rightarrow x_{Y=1} = \sqrt{u} \quad (11)$$

```
In [2]: def sample_of_x(y):
```

```
    """
```

```
    This function applies the inverse transform sampling. The parameter Y can be 0 or 1.
```

```
    """
```

```
    u = np.random.rand(*np.empty_like(y).shape)
```

```
    x = np.empty_like(y, dtype=float)
```

```

x[y == 0] = 1 - np.sqrt(1 - u[y == 0])
x[y == 1] = np.sqrt(u[y == 1])

return x

def create_data(N):
    """
    This function creates a set of N labels and the corresponding features.
    """
    # draw a uniform variable, set y to 0, if below 0.5, 1 if larger than 0.5.
    # =>  $p(y = 0) = p(y = 1) = 0.5$ 
    labels = np.array([0 if np.random.rand() < 0.5 else 1 for _ in range(N)])
    features = sample_of_x(labels)
    return features, labels

```


In [3]: # Checking the distributions

```

features, labels = create_data(100000)

fig, ax = plt.subplots(1, 1)
ax.hist(features[labels == 0], normed=1, bins=60, alpha=0.7, label="$Y = 0$")
ax.hist(features[labels == 1], normed=1, bins=60, alpha=0.7, label="$Y = 1$")
ax.set_xlabel('x')
ax.set_ylabel('p(x)')
ax.legend()
plt.show()

```



You can see in the plot, that for $Y = 0$ the distribution follows $p(x) = 2x - 2$ and for $Y = 1$ the distribution follows $p(x) = 2x$.

2 1.2 Classification by threshold

```
In [4]: def thresh_rule_a(x, t):
        """
        This function returns the result (label) of rule A for a given t and x.
        """
        f = np.empty_like(x, dtype=int)
        f[x <= t] = 0
        f[x > t] = 1
        return f

def thresh_rule_b(x, t):
    """
    This function returns the result (label) of rule B for a given t and x.
    """
    f = np.empty_like(x, dtype=int)
    f[x <= t] = 1
    f[x > t] = 0
    return f

def analytic_error_rule_a(t):
    """
    This function returns the analytic error of rule A for a given t
    """
    return 1/4 + (t - 1/2)**2

def analytic_error_rule_b(t):
    """
    This function returns the analytic error of rule B for a given t
    """
    return 3/4 - (t - 1/2)**2

def calculate_error(decisions, labels):
    """
    This function calculates the error rate of the decisions with the real labels.
    """
    diff = decisions - labels
    return np.count_nonzero(diff)/labels.shape[0]
```

2.1 Get error for different thresholds

```
In [5]: for t in [0.2, 0.5, 0.6]:
        features, labels = create_data(10000)

        decision_a = thresh_rule_a(features, t)
        decision_b = thresh_rule_b(features, t)

        print('t = {:.1f}:'.format(t))
        print('\t p_err_a = {:.3f} (analytic: {:.3f})'.format(calculate_error(decision_a, labels),
                                                                calculate_error_analytic(decision_a, labels)))
        print('\t p_err_b = {:.3f} (analytic: {:.3f})'.format(calculate_error(decision_b, labels),
                                                                calculate_error_analytic(decision_b, labels)))

t = 0.2:
    p_err_a = 0.345 (analytic: 0.340)
    p_err_b = 0.655 (analytic: 0.660)
t = 0.5:
    p_err_a = 0.255 (analytic: 0.250)
    p_err_b = 0.745 (analytic: 0.750)
t = 0.6:
    p_err_a = 0.259 (analytic: 0.260)
    p_err_b = 0.741 (analytic: 0.740)
```

As you can see above, the best results are reached for $t = 0.5$ with decision rule A (≈ 0.25). The observations fit very well with the analytic results.


Now we want to see, how test set size influence the mean value and standard deviation of p_{error} . Therefore, we calculate p_{error} for rule A with $t = 0.5$ 10 times to estimate the mean and standard-deviation. This is repeated for different test set sizes.

```
In [6]: means = []
        errors = []
        ms = [10, 100, 1000, 10000, 100000]
        for m in ms:
            test_results = []
            for _ in range(100):
                features, labels = create_data(m)
                decision_a = thresh_rule_a(features, 0.5)
                test_results.append(calculate_error(decision_a, labels))
            means.append(np.mean(test_results))
            errors.append(np.std(test_results))

In [7]: fig, ax = plt.subplots(1, 1)
        ax.plot(ms, means, label = "Mean")
        ax.plot(ms, errors, label = "Std")

        ax.set_xlabel('Test set size')
        ax.legend()

        plt.show()
```



As you can see in the plot, the mean value drops to 25% already for a small test set size. The standard deviation decreases towards 0. The decay is quite rapid, for 20000 elements it is already negligible against the error for 10 elements ($\approx 15\%$).

3 1.3 Nearest Neighbour classifier

In [8]: `class NearestNeighbour1D:`

```
def __init__(self):
    self.training = {"features": [], "labels": []}

def forget(self):
    self.training = {"features": [], "labels": []}

def train(self, features, labels):
    self.training = {"features": features, "labels": labels}

def classify(self, x):
    # get index of nearest neighbour
    ndx = np.argmin(np.abs(x - self.training["features"]))
    # return label of this guy
    return self.training["labels"][ndx]

def test(self, features, labels):
    errors = 0
```

```

    for i in range(len(features)):
        result = self.classify(features[i])
        if result != labels[i]:
            errors += 1
    return errors/len(features)

```

In [9]: *# create a classifier*

```

NN = NearestNeighbour1D()

```

```

error_rates = []

```

```

for _ in range(100):
    # reset this guy
    NN.forget()
    # create training set and train classifier
    training_features = [0, 0]
    training_labels = [0, 0]
    while training_labels[0] == training_labels[1]:
        training_features, training_labels = create_data(2)

    NN.train(training_features, training_labels)
    # create test set and test classifier
    test_features, test_labels = create_data(20000) # 20000 from above
    error_rates.append(NN.test(test_features, test_labels))

```

In [10]: `print('Average error rate: {:.f}, Standarddeviation: {:.f}'.format(np.mean(error_rates)`

```

Average error rate: 0.339008, Standarddeviation: 0.159542

```

In the range of the standard deviation, the average fits to the expected value of 35%.