

# generative-models


November 28, 2017

```
In [5]: import numpy as np
import h5py
import bisect
from matplotlib import pyplot as plt

In [6]: f = h5py.File('./digits.h5')
images = f["images"].value
labels = f["labels"].value
f.close()

images = images.reshape(len(images), 81)/255

In [7]: fig, ax = plt.subplots(2, 4)
for i in range(8):
    ax[i // 4, i % 4].imshow(images[labels == 3][i].reshape(9, 9), cmap='Greys')
plt.show()
```



# 1 Question 1: Sampling from Naive Bayes

This is the working example provided by Prof. Köthe.

```
In [8]: class Histogram(object):
        """Histogram.
        """

        def __init__(self, values):
            """Create a histogram for the given values using the Freedman-Diaconis rule for
            the number of bins.

            :param values: the values from which the histogram will be computed
            """
            self.num_instances = len(values)

            # Get the number of bins.
            v_min, v_25, v_75, v_max = np.percentile(values, [0, 25, 75, 100])
            # freedman_diaconis_width = 2 * (v_75 - v_25) / (len(values) ** (1/3.0))
            freedman_diaconis_width = (v_max - v_min) / (len(values) ** (1/3.0))
            num_bins = int(round((v_max - v_min) / freedman_diaconis_width))
            assert num_bins > 0

            # Fill the bins.
            self.heights, self.bin_edges = np.histogram(values, bins=num_bins)

        def find_bin(self, value):
            """Find the bin index of the given value.

            :param value: some value
            :return: bin index
            """
            bin_index = bisect.bisect_left(self.bin_edges, value) - 1
            bin_index = max(bin_index, 0)
            bin_index = min(bin_index, len(self.heights)-1)
            return bin_index

        def bin_probability(self, bin_index):
            """Return the bin probability of the desired bin.

            :param bin_index: index of the bin
            :return: probability of the bin
            """
            assert 0 <= bin_index <= len(self.heights) - 1
            return self.heights[bin_index] / float(self.num_instances)

class NaiveBayesClassifier(object):
    """Naive Bayes classifier.
    """
```

```

def __init__(self):
    self.num_instances = None # Number of instances that were used in training.
    self.num_feats = None # Number of features that were used in training.
    self.classes = None # The classes that were found in training.
    self.histograms = {} # Dict with histograms, key: class, value: list with his
    self.priors = {} # Dict with priors, key: class, value: prior of the class.

def train(self, train_x, train_y):
    """Train the classifier.

    :param train_x: training x data
    :param train_y: training y data
    """

    assert train_x.shape[0] == len(train_y)
    self.num_instances = train_x.shape[0]
    self.num_feats = train_x.shape[1]
    self.classes = np.unique(train_y)

    # Create one histogram for each class and each feature.
    self.histograms.clear()
    self.priors.clear()
    for cl in self.classes:
        # Get the data of the current class.
        train_x_cl = [train_x[k] for k in range(self.num_instances) if train_y[k] == cl]
        num_instances_cl = len(train_x_cl)
        self.priors[cl] = num_instances_cl / float(self.num_instances)

        # Create one histogram per feature.
        self.histograms[cl] = []
        for i in range(self.num_feats):
            histo_points = [train_x_cl[k][i] for k in range(num_instances_cl)]
            self.histograms[cl].append(Histogram(histo_points))

def probabilities_single(self, test_x):
    """Compute the class probabilities of a single instance.

    :param test_x: test x data of a single instance
    :return: probabilities of test_x for each class
    """

    assert len(test_x) == self.num_feats, \
        "NaiveBayesClassifier.probabilities_single(): Number of features in test and train data does not match."

    probs = {}
    for cl in self.classes:
        # Compute the likelihood.
        prob = self.priors[cl]
        for i in range(self.num_feats):
            # Find height of histogram bin with class cl that contains test_x[i].

```

```

        histo = self.histograms[c1][i]
        bin_index = histo.find_bin(test_x[i])
        prob *= histo.bin_probability(bin_index)
        probs[c1] = prob
    return probs

def predict_single(self, test_x):
    """Predict the class of a single instance.

    :param test_x: test x data of a single instance
    :return: predicted class of test_x
    """
    probs = self.probabilities_single(test_x)
    max_cl = max(probs.keys(), key=lambda cl: probs[cl])
    return max_cl

def predict(self, test_x):
    """Predict the classes of the given sample.

    :param test_x: test x data
    :return: predicted classes of test_x
    """
    assert test_x.shape[1] == self.num_feats, \
        "NaiveBayesClassifier.predict(): Number of features in test and training mismatch"
    return np.array([self.predict_single(test_x[i]) for i in range(test_x.shape[0])])

```

```

In [9]: nb = NaiveBayesClassifier()
        nb.train(images[labels == 3], labels[labels == 3])

```

## 1.1 Sampling from histograms

```

In [10]: def first_above_zero(z):
        z[z < 0] = 2 # disable all values where q is smaller than t, but do not delete them
        return np.argmin(z) # get the index

def sample_naive_bayes(histograms):
    q = [np.cumsum(histogram.heights)/np.sum(histogram.heights) for histogram in histograms]

    # Print histogram and cumulative
    #fig, ax = plt.subplots(2, 1)
    #i = 27
    #width = histograms[i].bin_edges[1] - histograms[i].bin_edges[0]
    #ax[0].bar(histograms[i].bin_edges[:-1] + 0.5*width, histograms[i].heights, width)
    #ax[1].bar(histograms[i].bin_edges[:-1] + 0.5*width, q[i], width)
    #plt.show()

    t = np.random.rand(len(histograms))

```

```

l = [first_above_zero(cum - t[i]) for i, cum in enumerate(q)]
x = np.random.rand(len(histograms))


return np.array([histogram.bin_edges[l[i]] + x[i]*(histogram.bin_edges[1] - histo

```

```

In [11]: fig, ax = plt.subplots(2, 4)
for i in range(8):
    ax[i // 4, i % 4].imshow(sample_naive_bayes(nb.histograms[3]).reshape(9, 9), cmap=
plt.show()

```



## 2 Question 2: Density trees

$n$  is the number of cuts we make,  $i$  is the number of the perfect split. Maximizing score  $s$

$$s = -p_l^2 * V_l + p_\lambda^2 * V_\lambda + p_\rho^2 * V_\rho \quad (1)$$

$$= -\frac{p_l^2}{V_l} + \frac{p_\lambda^2}{V_\lambda} + \frac{p_\rho^2}{V_\rho} \quad (2)$$

$$= -\frac{p_l^2}{V_l} + \frac{p_\lambda^2}{i * V_l / n} + \frac{p_\rho^2}{(n - i) * V_l / n} \quad (3)$$

$$= -\frac{p_l^2}{V_l} + \frac{n}{V_l} * \left( \frac{p_\lambda^2}{i} + \frac{p_\rho^2}{n - i} \right) \quad (4)$$

Here we used, that

$$V_\rho = d_{\rho,1} * \dots * d_{\rho,D} = \frac{i}{n} * d_{\lambda,1} * \dots * d_{\lambda,D} = \frac{i}{n} * V_\lambda \quad (5)$$

```

In [87]: def volume(a):
        """
        Calculates the n-dimensional volume.
        """
        return np.prod(a[:,1])

class Node:
    def __init__(self, data, bbox, vol, tot_n):
        self.left = None
        self.right = None

        self.data = data # the actual data
        self.bbox = bbox # boundary box of the data, 2 dimensions, first is the start

        self.vol = vol # the calculated volume
        self.prob = len(data)/tot_n # probability of the node
        self.tot_n = tot_n # yeah it is crep to put it in every node, but it is late.

        self.best_split = None

    @staticmethod
    def score_min_error(n_l, n_r, tot_n, vol, i, n):
        # NOTE: maximizing this could be done with much less calculations
        return -(((n_l+n_r)/tot_n)**2)/vol) + n*(((n_l/tot_n)**2)/i) + (((n_r/tot_n)**2)/i)

    @staticmethod
    def score_max_uniform(n_l, n_r, tot_n, vol, i, n):
        vol_left = vol*i/n
        return ((n_l*vol - (n_r + n_l)*vol_left)**2)/(vol_left*(vol - vol_left))

    def get_score(self, score_func, n = 10):

        # if has been already calculated
        if self.best_split != None:
            return self.best_split["score"]

        # init
        self.best_split = {
            "score": 0,
            "left": None,
            "right": None,
            "dimension": None,
            "ratio": None
        }

        for d in range(self.data[0].shape[0]): # loop over all dimensions
            for i in range(1, n): # loop over the number of cuts

```

```

        # split the data into left and right box
        data_left = self.data[self.data[:,d] < self.bbox[d,0] + i*self.bbox[d,1]-self.bbox[d,0]]
        data_right = self.data[self.data[:,d] >= self.bbox[d,0] + i*self.bbox[d,1]-self.bbox[d,0]]

        # do not allow split which would just shrink the node
        # TODO: right?
        if len(data_left) == 0 or len(data_right) == 0:
            continue

        # assert len(data_left) + len(data_right) == len(self.data)

        # call the given scoring function
        score = score_func(len(data_left), len(data_right), self.tot_n, self.vol)

        # if found a new highest score keep the current values
        if score > self.best_split["score"]:
            self.best_split["score"] = score
            self.best_split["left"] = data_left
            self.best_split["right"] = data_right
            self.best_split["dimension"] = d
            self.best_split["ratio"] = i/n

    return self.best_split["score"]

def split(self):

    # assert len(self.data) != 0

    bbox_left = np.copy(self.bbox)
    bbox_left[self.best_split["dimension"], 1] *= self.best_split["ratio"] # shrink
    self.left = Node(self.best_split["left"], bbox_left, self.vol*self.best_split["ratio"])

    bbox_right = np.copy(self.bbox)
    bbox_right[self.best_split["dimension"], 0] += bbox_right[self.best_split["dimension"], 1]*self.best_split["ratio"]
    bbox_right[self.best_split["dimension"], 1] *= (1 - self.best_split["ratio"])
    self.right = Node(self.best_split["right"], bbox_right, self.vol*(1 - self.best_split["ratio"]))

    ##### Tests on bounding boxes
    # Feel free to test but should pass :D
    # assert (volume(bbox_left) - self.left.vol) < 1e-6
    # assert (volume(bbox_right) - self.right.vol) < 1e-6
    # for i in range(81):
    #     if i == self.best_split["dimension"]:
    #         continue
    #     assert self.left.bbox[i, 0] == self.right.bbox[i, 0]
    #     assert self.left.bbox[i, 1] == self.right.bbox[i, 1]
    # assert self.right.bbox[self.best_split["dimension"], 0] == self.left.bbox[self.best_split["dimension"], 0]

```

```

@staticmethod
def get_leaf_of_instance(node, instance):
    current = node
    while current.left != None: # avoid recursion
        if instance[current.best_split["dimension"]] > current.right.bbox[current
            current = current.right
        else:
            current = current.left
    return current

```

## 2.1 Training

In [90]: # DEBUG

```

def get_probability(node):
    """
    Recursively check the probability of the node and compare to precomputed one.
    """
    if node.left == None:
        return node.prob
    else:
        a = get_probability(node.left)
        b = get_probability(node.right)
        if np.abs(a + b - node.prob) > 1e-6:
            print(a, b, node.prob)
            raise ValueError('FATAL ERROR: probabilities do not match')
        return a + b

# DEBUG
def check_instances(node):
    """
    Check if each datapoint is inside the given boundary box
    """
    for i in node.data:
        for j in range(node.data[0].shape[0]):
            # should both be 0 for both
            if node.bbox[j, 0] - i[j] > 1e-6 or i[j] - node.bbox[j, 0] > 1e-6:
                print(i[j], node.bbox[j])
                raise ValueError('FATAL ERROR: instances outside of bounding box')

def fit_density_tree1(features, tau, score_func):
    bincount = int(tau*len(features)*(1/3))
    print('Use {:d} bins'.format(bincount))

    root_bounding = np.empty((features[0].shape[0], 2)) # init boundary box

```



```

for i in range(features[0].shape[0]): # for each feature...
    root_bounding[i, 0] = np.min(features[:,i]) # ... start box at smallest value
    root_bounding[i, 1] = np.max(features[:,i]) - np.min(features[:,i]) # ... and

# init root node
root = Node(features, root_bounding, volume(root_bounding), len(features))

# init leafs array
leafs = [root]

while len(leafs) < bincount:
    if int(bincount/100) == 0 or len(leafs) % int(bincount/100):
        print('Leafs: {:d}, {:.0f}%'.format(len(leafs), 100*len(leafs)/bincount),

        # get the scores for each leaf
        scores = np.array([node.get_score(score_func) for node in leafs])

        # get the leaf with the highest score
        leaf_index = np.argmax(scores)

        # if the highest score is below 0: abort
        if scores[leaf_index] <= 0:
            print('Highest score below 0, do not split any further.')
            break

        # split the leaf with highest score
        leafs[leaf_index].split()

        # append its children to the leaf
        leafs.append(leafs[leaf_index].left)
        leafs.append(leafs[leaf_index].right)

        # remove from leafs
        del leafs[leaf_index]

print('Finished.', ' '*30)
#print('Check for errors...')
# DEBUG
#for i in leafs:
#    check_instances(i)
#print('All instances are inside their boundaries in the leaf nodes')

# DEBUG
#print('Root node has probability of {:.f}'.format(get_probability(root)))
return root

```

In [91]: tree\_min\_error = fit\_density\_tree1(images[labels == 3], 10, Node.score\_min\_error)

Use 230 bins

Finished.

```
In [92]: tree_max_unif = fit_density_tree1(images[labels == 3], 10, Node.score_max_uniform)
```

Use 230 bins

Leafs: 151, 66%

/Users/maxsimon/anaconda/lib/python3.6/site-packages/ipykernel\_launcher.py:29: RuntimeWarning:

Leafs: 153, 67%

/Users/maxsimon/anaconda/lib/python3.6/site-packages/ipykernel\_launcher.py:29: RuntimeWarning:


Finished.9, 100%

## 2.2 Sampling

```
In [105]: def sample_density_tree(tree):
            t = np.random.rand()
            current = tree
            while current.left != None:
                if t <= current.prob:
                    current = current.left
                else:
                    t -= current.prob
                    current = current.right
            # current is now a leaf
            x = np.random.rand(current.data[0].shape[0]) # random numbers for inside the bin
            return current.bbox[:,0] + x*current.bbox[:,1]
```


### 2.2.1 Minimizing error

```
In [106]: fig, ax = plt.subplots(2, 4)
            for i in range(8):
                ax[i // 4, i % 4].imshow(sample_density_tree(tree_min_error).reshape(9, 9), cmap=
            plt.show()
```



## 2.2.2 Maximizing uniformity

```
In [107]: fig, ax = plt.subplots(2, 4)
          for i in range(8):
              ax[i // 4, i % 4].imshow(sample_density_tree(tree_max_unif).reshape(9, 9), cmap=
plt.show()
```



## 2.3 Classification

```
In [97]: def predict_density_trees(test_features, tree_for_3, tree_for_9):
        predicted_labels = np.empty(len(test_features))
        for i, feature in enumerate(test_features):
            if Node.get_leaf_of_instance(tree_for_3, feature).prob < Node.get_leaf_of_instance(tree_for_9, feature).prob:
                predicted_labels[i] = 9
            else:
                predicted_labels[i] = 3
        return predicted_labels
```

```
In [98]: #tree_3 = fit_density_tree1(images[labels == 3], 10, Node.score_min_error) already calculated
        tree_9 = fit_density_tree1(images[labels == 9], 10, Node.score_min_error)
```

Use 228 bins

Finished.


```
In [99]: def get_confusion_matrix(predicted, truth, possible_labels=[3, 9]):
        conf = np.empty((len(possible_labels), len(possible_labels)))
        for i, k in enumerate(possible_labels):
            items, counts = np.unique(predicted[truth == k], return_counts=True)
            count_array = np.array([counts[np.where(items == tested_k)[0][0]] for tested_k in possible_labels])
            conf[i] = count_array
        return conf
```

```
In [100]: f = h5py.File('./digits_test.h5')
        images_test = f["images"].value
        labels_test = f["labels"].value
        f.close()
        images_test = images_test.reshape(len(images_test), 81)/255
```

```
In [101]: mask = np.logical_or(labels_test == 3, labels_test == 9)
        prediction = predict_density_trees(images_test[mask], tree_min_error, tree_9)
        error_rate = np.count_nonzero(prediction - labels_test[mask])/prediction.shape[0]
        conf_ma = get_confusion_matrix(prediction, labels_test[mask])
```

```
In [108]: print('Error rate: {:.5f}'.format(error_rate))
        fig, ax = plt.subplots(1, 1)
        cax = ax.imshow(conf_ma)
        plt.colorbar(cax)
        plt.show()
```

Error rate: 0.39222



Yep, the 3s look not like a typical 3 and the error rate is enormous. The problem is, that we get very large bins

### 3 Question 3: Sampling with QDA

```
In [109]: def get_class(x, y, desired):
           """
           Returns a subarray of x where y = desired
           """
           return x[y == desired]

def fit_qda(training_features, training_labels, possible_features=[3]):
    D = training_features.shape[1] # dimension of features
    F = 1

    # create an array of testsetdata
    ts = [get_class(training_features, training_labels, feature) for feature in possible_features]
    # calculate the total amount of testdata (including all possible features)
    N_tot = sum([tsf.shape[0] for tsf in ts])

    # mu has the shape FxD
    mu = np.empty((F, D), dtype=float)
    # cov should have the shape FxDxD
    cov = np.empty((F, D, D))
    # the priors are scalars and have therefore the shape F
    p = np.empty(F)
```

```

for i in range(F):
    N = ts[i].shape[0] # number of training instances for the feature possible_f
    # calculating mu
    mu[i] = np.mean(ts[i], axis=0)

    # calculating the covariance matrix
    ts_centralised = ts[i] - mu[i]
    # some numpy magic
    cov[i] = np.add.reduce(ts_centralised[:, :, np.newaxis] * ts_centralised[:, np.newaxis, :], axis=0)

    # calculating the priors
    p[i] = N/N_tot


# done
return mu, cov, p

```

```
In [110]: mu, cov, p = fit_qda(images[labels == 3], labels[labels == 3])
```

```
In [111]: def sample_qda(mu, cov):
    return np.random.multivariate_normal(mu, cov)
```

```
In [112]: fig, ax = plt.subplots(2, 4)
    for i in range(8):
        ax[i // 4, i % 4].imshow(sample_qda(mu[0], cov[0]).reshape(9, 9), cmap='Greys')
    plt.show()
```



```
In [ ]:
```